



**RAJALAKSHMI ENGINEERING COLLEGE**

*Approved by AICTE | Affiliated to Anna University | Accredited by NAAC*

**Department of Computer Science and Engineering**

**CS23334 Fundamentals of Data Science Lab**

**III semester II Year (2023R)**

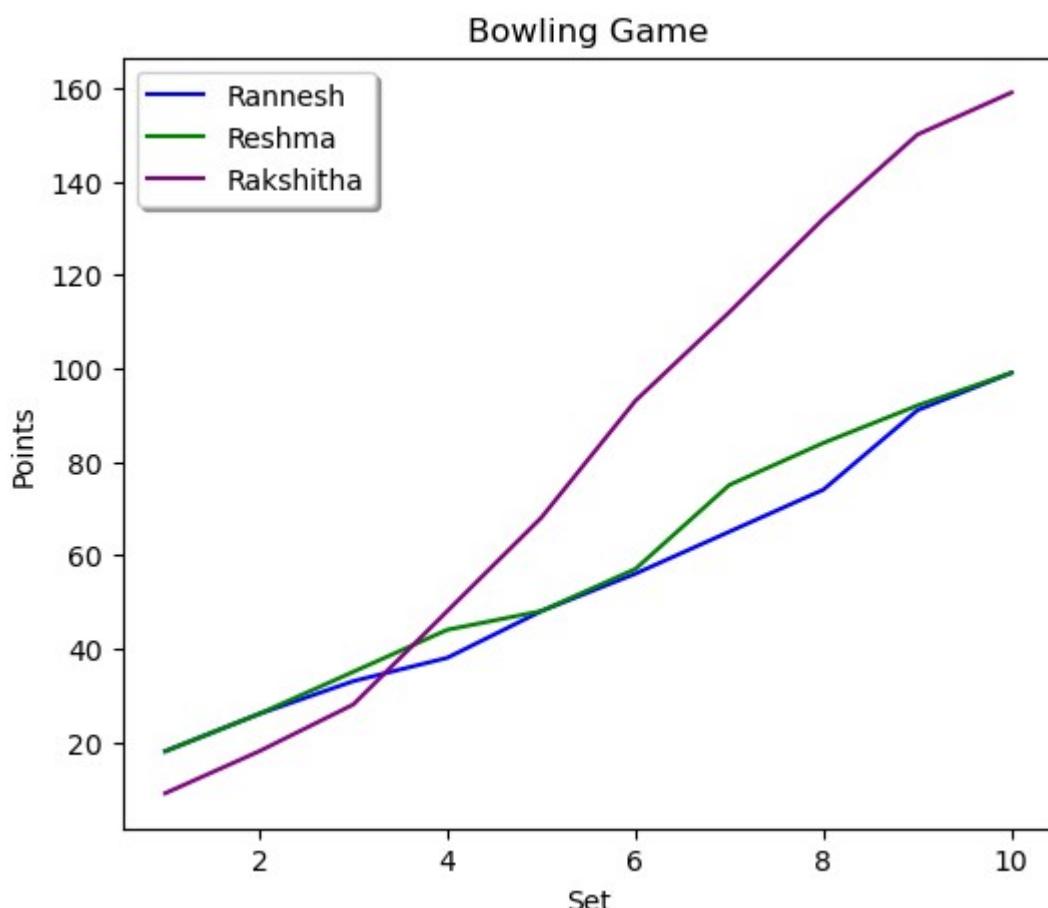
**Name of the Student : RANNESH KHUMAR B R**

**Register Number : 240701422**



```
#Rannesh Khumar B R  
#240701422  
#Fundamentals of Data Science  
#17.07.2025  
#LinePlot,Bargraph,Piechart,Histogram and Scatter plot using matplotlib
```

```
import matplotlib.pyplot as plt  
set=list(range(1,11))  
p1=[18,26,33,38,48,56,65,74,91,99]  
p2=[18,26,35,44,48,57,75,84,92,99]  
p3=[9,18,28,48,68,93,112,132,150,159]  
plt.figure(figsize=(6,5))  
plt.plot(set,p1,color='blue',label='Rannesh')  
plt.plot(set,p2,color='green',label='Reshma')  
plt.plot(set,p3,color='purple',label='Rakshitha')  
plt.title("Bowling Game")  
plt.xlabel("Set")  
plt.ylabel("Points")  
plt.legend(loc='upper left',shadow=True)  
plt.show()
```

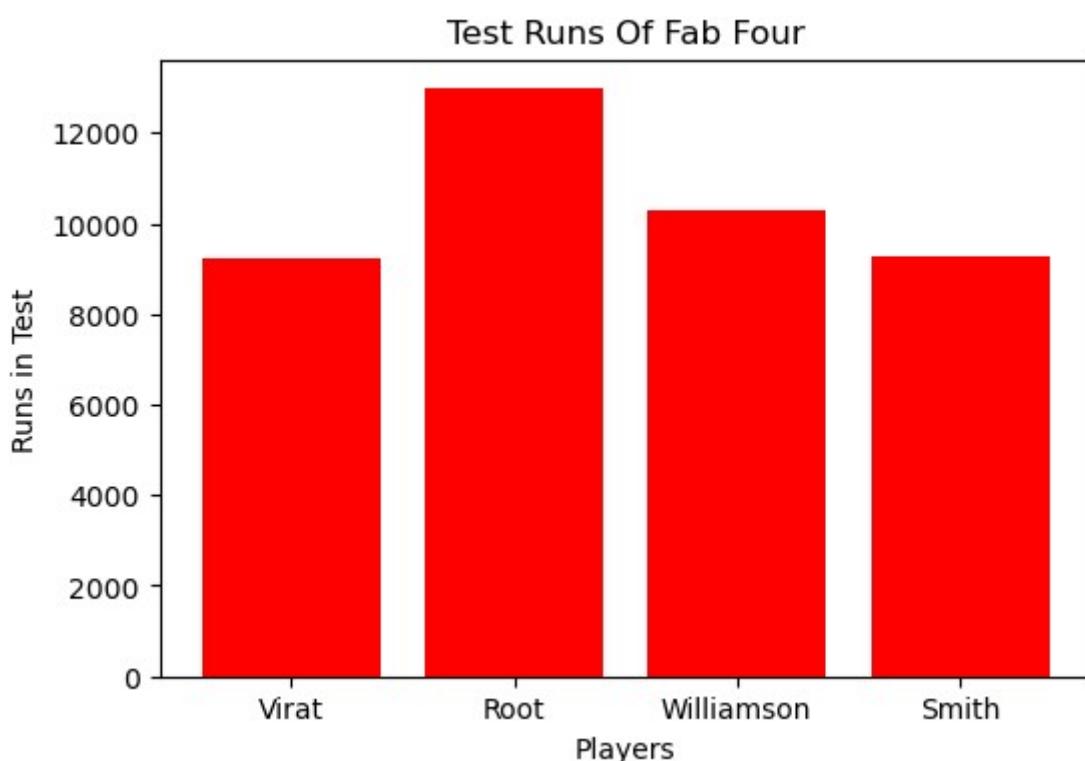


```
import matplotlib.pyplot as plt

categories = ['Virat', 'Root', 'Williamson', 'Smith']

values = [9230, 12972, 10271, 9276]

plt.figure(figsize=(6, 4))
plt.bar(categories, values, color='red')
plt.title("Test Runs Of Fab Four")
plt.xlabel("Players")
plt.ylabel("Runs in Test")
plt.show()
```

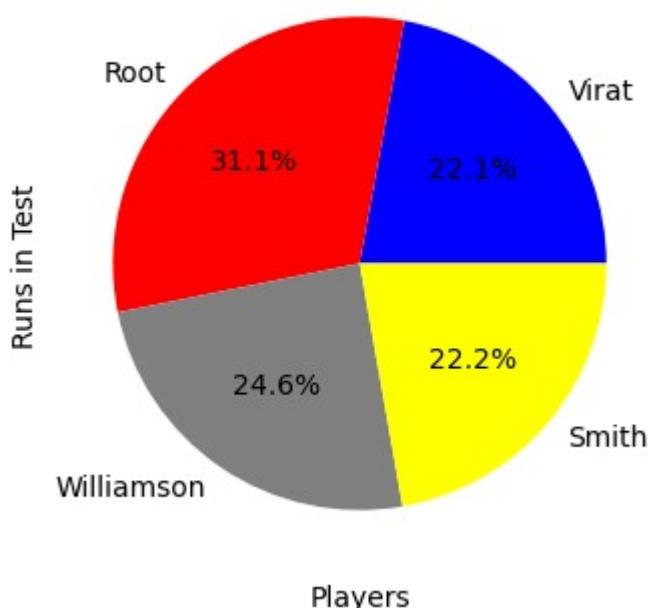


```
import matplotlib.pyplot as plt

categories = ['Virat', 'Root', 'Williamson', 'Smith']

values = [9230, 12972, 10271, 9276]
color=['blue','red','grey','yellow']
plt.figure(figsize=(6, 4))
plt.pie(values, labels=categories, colors=color, autopct='%1.1f%%')
plt.title("Test Runs Of Fab Four")
plt.xlabel("Players")
plt.ylabel("Runs in Test")
plt.show()
```

**Test Runs Of Fab Four**



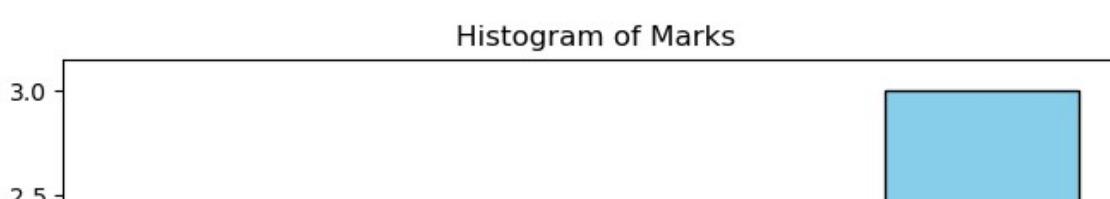
```
import matplotlib.pyplot as plt

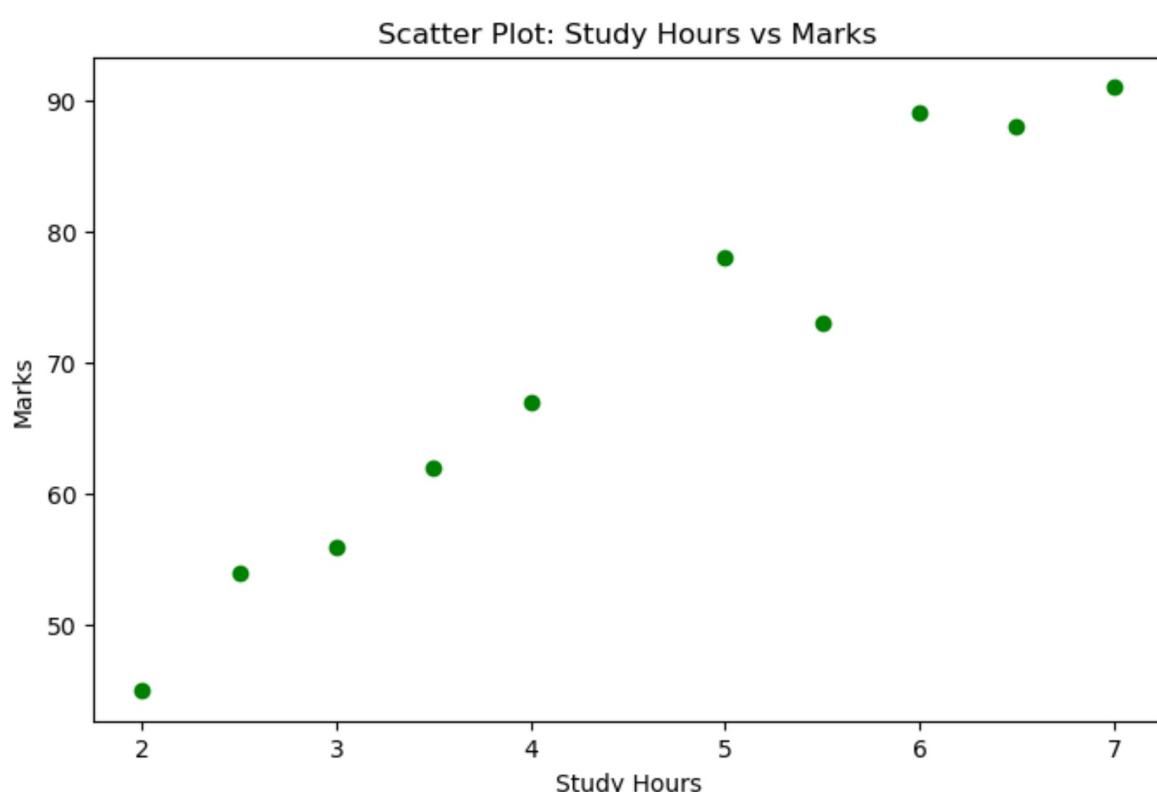
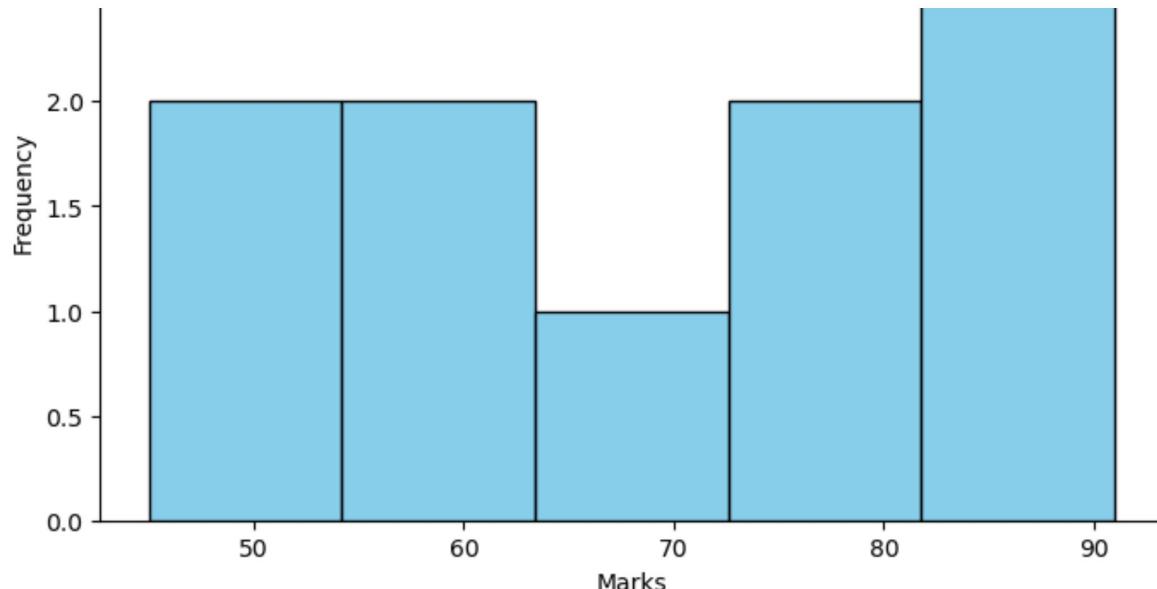
# Manually entered data
marks = [45, 78, 89, 56, 67, 91, 62, 73, 88, 54]
study_hours = [2, 5, 6, 3, 4, 7, 3.5, 5.5, 6.5, 2.5]

# Histogram for Marks
plt.figure(figsize=(8, 5))
plt.hist(marks, bins=5, color='skyblue', edgecolor='black')
plt.title('Histogram of Marks')
plt.xlabel('Marks')
plt.ylabel('Frequency')

plt.show()

# Scatter Plot of Study Hours vs Marks
plt.figure(figsize=(8, 5))
plt.scatter(study_hours, marks, color='green', marker='o')
plt.title('Scatter Plot: Study Hours vs Marks')
plt.xlabel('Study Hours')
plt.ylabel('Marks')
plt.show()
```





Start coding or generate with AI.





```
# Rannesh Khumar B R  
# 240701422  
# 24.7.25  
# Data preprocessing
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
file_path='sales_data(2).csv'  
df = pd.read_csv(file_path)
```

```
df
```

	Date	Product	Sales	Quantity	Region
0	01-01-2023	Product A	200	4	North
1	02-01-2023	Product B	150	3	South
2	03-01-2023	Product A	220	5	North
3	04-01-2023	Product C	300	6	East
4	05-01-2023	Product B	180	4	West
5	06-01-2023	Product A	210	5	North
6	07-01-2023	Product C	320	7	East
7	08-01-2023	Product B	160	3	South
8	09-01-2023	Product A	230	6	North
9	10-01-2023	Product C	310	7	East
10	11-01-2023	Product B	190	4	West
11	12-01-2023	Product A	240	6	North
12	13-01-2023	Product C	330	8	East
13	14-01-2023	Product B	170	3	South
14	15-01-2023	Product A	250	7	North
15	16-01-2023	Product C	340	8	East

```
df['Sales'].fillna(df['Sales'].mean())  
df.dropna(subset=['Product', 'Quantity', 'Region'])
```

	Date	Product	Sales	Quantity	Region
0	01-01-2023	Product A	200	4	North
1	02-01-2023	Product B	150	3	South
2	03-01-2023	Product A	220	5	North
3	04-01-2023	Product C	300	6	East
4	05-01-2023	Product B	180	4	West
5	06-01-2023	Product A	210	5	North
6	07-01-2023	Product C	320	7	East
7	08-01-2023	Product B	160	3	South
8	09-01-2023	Product A	230	6	North
9	10-01-2023	Product C	310	7	East
10	11-01-2023	Product B	190	4	West
11	12-01-2023	Product A	240	6	North
12	13-01-2023	Product C	330	8	East
13	14-01-2023	Product B	170	3	South
14	15-01-2023	Product A	250	7	North
15	16-01-2023	Product C	340	8	East

```
df.describe()
```

	Sales	Quantity
count	16.000000	16.000000
mean	237.500000	5.375000
std	64.031242	1.746425
min	150.000000	3.000000
25%	187.500000	4.000000
50%	225.000000	5.500000
75%	302.500000	7.000000
max	340.000000	8.000000

```
product_summary = df.groupby('Product').agg({  
    'Sales': 'sum',  
    'Quantity': 'sum'}
```

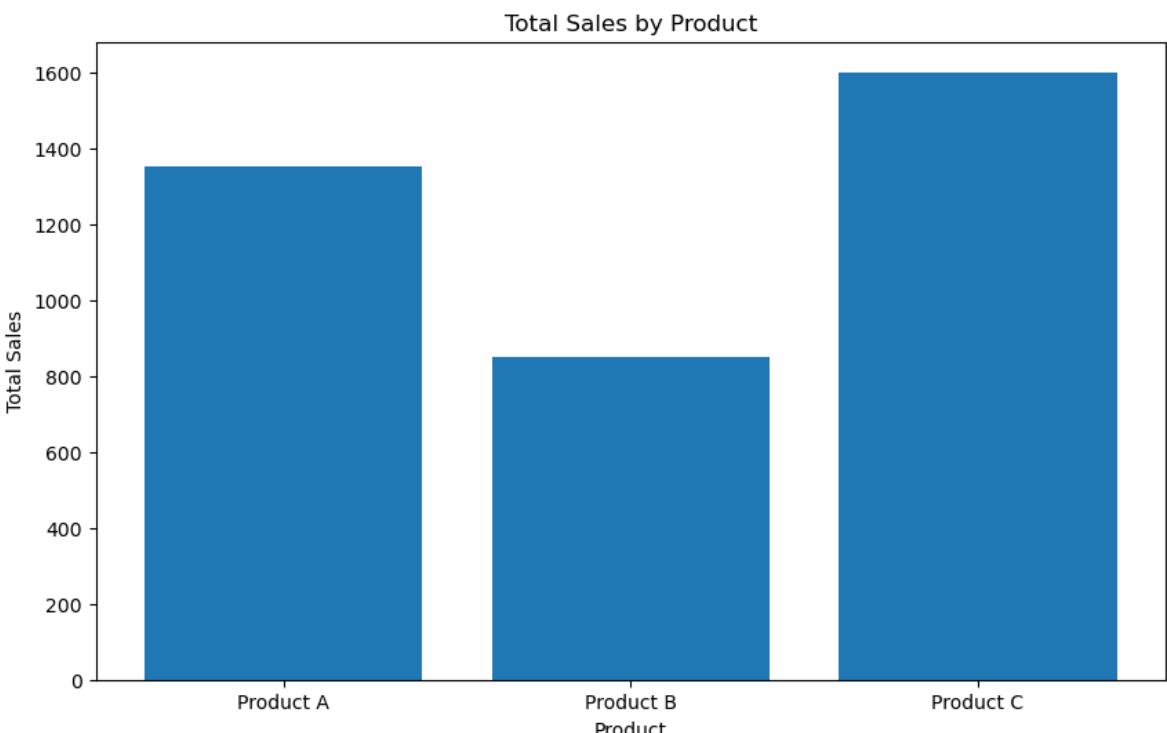
```
}).reset_index()
```

```
product_summary
```

	Product	Sales	Quantity
0	Product A	1350	33
1	Product B	850	17
2	Product C	1600	36

```
plt.figure(figsize=(10, 6))
plt.bar(product_summary['Product'], product_summary['Sales'])
plt.xlabel('Product')
plt.ylabel('Total Sales')
plt.title('Total Sales by Product')
```

```
Text(0.5, 1.0, 'Total Sales by Product')
```



```
sales_over_time = df.groupby('Date').agg({'Sales': 'sum'}).reset_index()
plt.figure(figsize=(10, 6))
plt.plot(sales_over_time['Date'], sales_over_time['Sales'])
plt.xlabel('Date')
plt.ylabel('Total Sales')
plt.title('SalesOver Time')
```

Text(0.5, 1.0, 'SalesOver Time')



```
pivot_table = df.pivot_table(values='Sales', index='Region', columns='Produc
```

pivot\_table

Region	Product A	Product B	Product C
--------	-----------	-----------	-----------

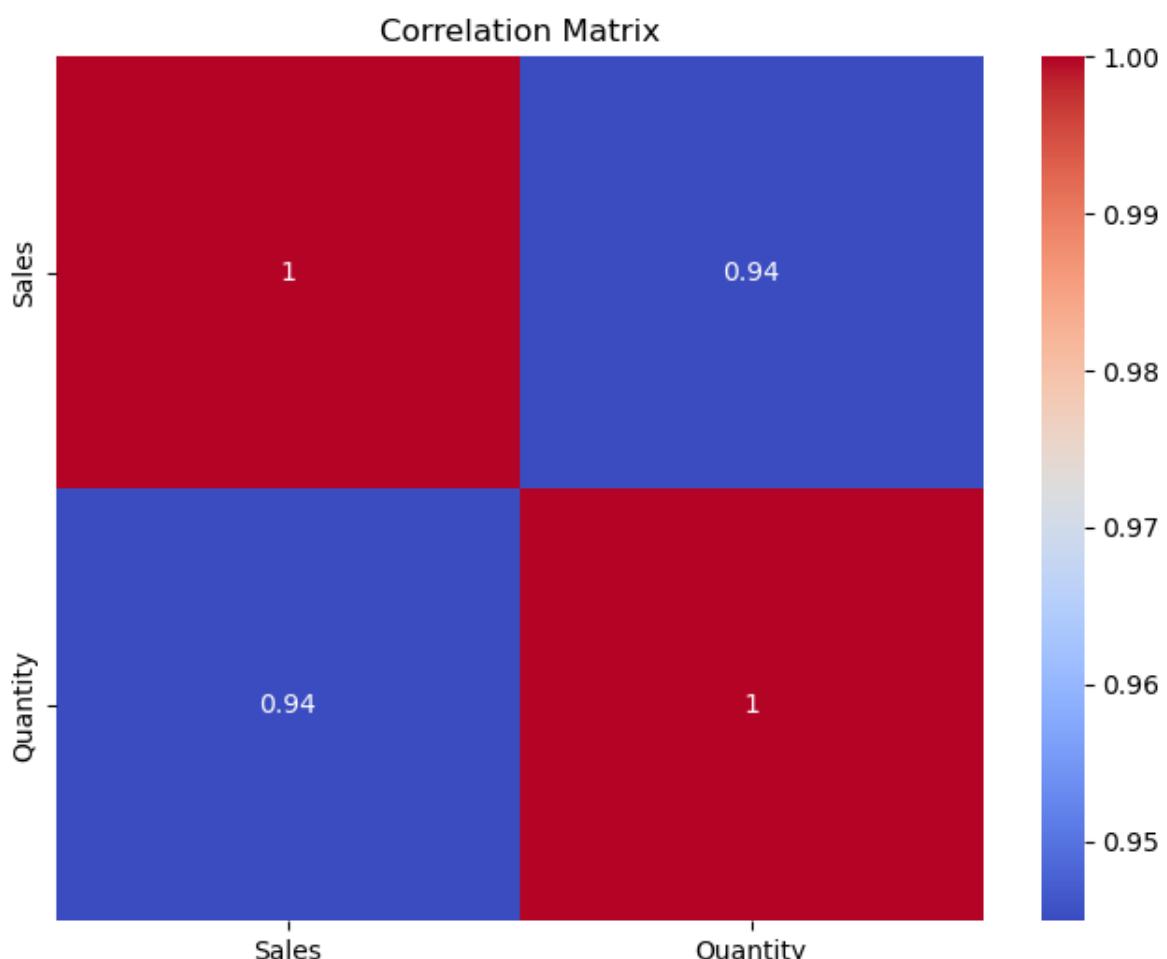
Region			
--------	--	--	--

East	0	0	1600
North	1350	0	0

<b>South</b>	0	480	0
<b>West</b>	0	370	0

```
correlation_matrix = df.corr()
```

```
import seaborn as sns
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



Start coding or generate with AI.





```
# Rannesh Khumar B R
# 240701422
# 7.8.25
# Data preprocessing
```

```
import numpy as np
import pandas as pd
df=pd.read_csv("pre_process_datasample(1).csv")
df
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes
5	France	35.0	58000.0	Yes
6	Spain	NaN	52000.0	No
7	France	48.0	79000.0	Yes
8	Germany	50.0	83000.0	No
9	France	37.0	67000.0	Yes

```
df.info()
df.Country.mode()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Country     10 non-null    object 
 1   Age         9 non-null    float64
 2   Salary       9 non-null    float64
 3   Purchased    10 non-null   object 
dtypes: float64(2), object(2)
memory usage: 452.0+ bytes
0   France
Name: Country, dtype: object
```

```
df.Country.mode()[0]
```

```
'France'
```

```
type(df.Country.mode())
```

```
pandas.core.series.Series
```

```
df['Country'] = df['Country'].fillna(df['Country'].mode()[0])
df['Age'] = df['Age'].fillna(df['Age'].median())
df['Salary'] = df['Salary'].fillna(round(df['Salary'].mean()))
```

```
pd.get_dummies(df.Country)
```

	France	Germany	Spain
0	True	False	False
1	False	False	True
2	False	True	False
3	False	False	True
4	False	True	False
5	True	False	False
6	False	False	True
7	True	False	False
8	False	True	False
9	True	False	False

```
updated_dataset=pd.concat([pd.get_dummies(df.Country),df.iloc[:,[1,2,3]]],axis=1)
updated_dataset
```

	France	Germany	Spain	Age	Salary	Purchased
0	True	False	False	44.0	72000.0	No
1	False	False	True	27.0	48000.0	Yes
2	False	True	False	30.0	54000.0	No
3	False	False	True	38.0	61000.0	No
4	False	True	False	40.0	63778.0	Yes
-	-	-	-	-	-	-

5	True	False	False	35.0	58000.0	Yes
6	False	False	True	38.0	52000.0	No
7	True	False	False	48.0	79000.0	Yes
8	False	True	False	50.0	83000.0	No
9	True	False	False	37.0	67000.0	Yes

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Country     10 non-null    object 
 1   Age         10 non-null    float64 
 2   Salary       10 non-null    float64 
 3   Purchased   10 non-null    object 
dtypes: float64(2), object(2)
memory usage: 452.0+ bytes
```

```
updated_dataset.Purchased.replace(['No','Yes'],[0,1])
```

```
0    0
1    1
2    0
3    0
4    1
5    1
6    0
7    1
8    0
9    1
Name: Purchased, dtype: int64
```

```
updated_dataset
```

	France	Germany	Spain	Age	Salary	Purchased
0	True	False	False	44.0	72000.0	0
1	False	False	True	27.0	48000.0	1
2	False	True	False	30.0	54000.0	0
3	False	False	True	38.0	61000.0	0
4	False	True	False	40.0	63778.0	1
5	True	False	False	35.0	58000.0	1
6	-	-	-	~	~	~

6	False	False	True	38.0	52000.0	0
7	True	False	False	48.0	79000.0	1
8	False	True	False	50.0	83000.0	0
9	True	False	False	37.0	67000.0	1

Start coding or [generate](#) with AI.



```
# Rannesh Khumar B R  
# 240701422  
# 7.8.25  
# Handling Missing Values
```

```
import numpy as np  
import pandas as pd  
df=pd.read_csv("Hotel_Dataset.csv")  
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPeople
0	1	20-25	4	Ibis	veg	1300	
1	2	30-35	5	LemonTree	Non-Veg	2000	
2	3	25-30	6	RedFox	Veg	1322	
3	4	20-25	-1	LemonTree	Veg	1234	
4	5	35+	3	Ibis	Vegetarian	989	
5	6	35+	3	Ibys	Non-Veg	1909	
6	7	35+	4	RedFox	Vegetarian	1000	
7	8	20-25	7	LemonTree	Veg	2999	
8	9	25-30	2	Ibis	Non-Veg	3456	
9	9	25-30	2	Ibis	Non-Veg	3456	
10	10	30-35	5	RedFox	non-Veg	-6755	

```
df.duplicated()
```

```
0    False  
1    False  
2    False  
3    False  
4    False  
5    False  
6    False  
7    False  
8    False  
9    True  
10   False  
dtype: bool
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 11 entries, 0 to 10
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CustomerID      11 non-null     int64  
 1   Age_Group       11 non-null     object  
 2   Rating(1-5)     11 non-null     int64  
 3   Hotel            11 non-null     object  
 4   FoodPreference   11 non-null     object  
 5   Bill             11 non-null     int64  
 6   NoOfPax          11 non-null     int64  
 7   EstimatedSalary  11 non-null     int64  
 8   Age_Group.1     11 non-null     object  
dtypes: int64(5), object(4)
memory usage: 924.0+ bytes

```

```

df.drop_duplicates(inplace=True)
df

```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax
0	1	20-25	4	Ibis	veg	1300	
1	2	30-35	5	LemonTree	Non-Veg	2000	
2	3	25-30	6	RedFox	Veg	1322	
3	4	20-25	-1	LemonTree	Veg	1234	
4	5	35+	3	Ibis	Vegetarian	989	
5	6	35+	3	Ibys	Non-Veg	1909	
6	7	35+	4	RedFox	Vegetarian	1000	
7	8	20-25	7	LemonTree	Veg	2999	
8	9	25-30	2	Ibis	Non-Veg	3456	
10	10	30-35	5	RedFox	non-Veg	-6755	

```

len(df)

```

```

index=np.array(list(range(0,len(df))))
df.set_index(index,inplace=True)
index
df

```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax
0	1	20-25	4	Ibis	veg	1300	
1	2	30-35	5	LemonTree	Non-Veg	2000	

2	3	25-30	6	RedFox	Veg	1322
3	4	20-25	-1	LemonTree	Veg	1234
4	5	35+	3	Ibis	Vegetarian	989
5	6	35+	3	Ibys	Non-Veg	1909
6	7	35+	4	RedFox	Vegetarian	1000
7	8	20-25	7	LemonTree	Veg	2999
8	9	25-30	2	Ibis	Non-Veg	3456
9	10	30-35	5	RedFox	non-Veg	-6755

```
df.drop(['Age_Group.1'],axis=1,inplace=True)
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfI
0	1	20-25	4	Ibis	veg	1300	
1	2	30-35	5	LemonTree	Non-Veg	2000	
2	3	25-30	6	RedFox	Veg	1322	
3	4	20-25	-1	LemonTree	Veg	1234	
4	5	35+	3	Ibis	Vegetarian	989	
5	6	35+	3	Ibys	Non-Veg	1909	
6	7	35+	4	RedFox	Vegetarian	1000	
7	8	20-25	7	LemonTree	Veg	2999	
8	9	25-30	2	Ibis	Non-Veg	3456	
9	10	30-35	5	RedFox	non-Veg	-6755	

```
df.CustomerID.loc[df.CustomerID<0]=np.nan
df.Bill.loc[df.Bill<0]=np.nan
df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
df
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_13572\2080958306.py:1: FutureWarning  
 You are setting values through chained assignment. Currently this works in ce  
 A typical example is when you are setting values in a column of a DataFrame,

```
df["col"][row_indexer] = value
```

Use `df.loc[row\_indexer, "col"] = values` instead, to perform the assignment

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.CustomerID.loc[df.CustomerID<0]=np.nan
C:\Users\HP\AppData\Local\Temp\ipykernel_13572\2080958306.py:1: SettingWithCo
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.CustomerID.loc[df.CustomerID<0]=np.nan
C:\Users\HP\AppData\Local\Temp\ipykernel_13572\2080958306.py:2: FutureWarning
You are setting values through chained assignment. Currently this works in ce
A typical example is when you are setting values in a column of a DataFrame,
```

```
df["col"][[row_indexer]] = value
```

Use `df.loc[row\_indexer, "col"] = values` instead, to perform the assignment

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.Bill.loc[df.Bill<0]=np.nan
C:\Users\HP\AppData\Local\Temp\ipykernel_13572\2080958306.py:2: SettingWithCo
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.Bill.loc[df.Bill<0]=np.nan
C:\Users\HP\AppData\Local\Temp\ipykernel_13572\2080958306.py:3: FutureWarning
You are setting values through chained assignment. Currently this works in ce
A typical example is when you are setting values in a column of a DataFrame,
```

```
df["col"][[row_indexer]] = value
```

Use `df.loc[row\_indexer, "col"] = values` instead, to perform the assignment

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
C:\Users\HP\AppData\Local\Temp\ipykernel_13572\2080958306.py:3: SettingWithCo
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
```

CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfReviews
0	1.0	20-25	4	Ibis	veg	1300.0
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0
2	3.0	25-30	6	RedFox	Veg	1322.0
3	4.0	20-25	-1	LemonTree	Veg	1234.0
4	5.0	35+	3	Ibis	Vegetarian	989.0
5	6.0	35+	3	Ibys	Non-Veg	1909.0
6	7.0	35+	4	RedFox	Vegetarian	1000.0
7	8.0	20-25	7	LemonTree	Veg	2000.0

1	8.0	20-25	1	LemonTree	veg	2999.0
8	9.0	25-30	2	Ibis	Non-Veg	3456.0
9	10.0	30-35	5	RedFox	non-Veg	NaN

```
df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan
df
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_13572\2129877948.py:1: FutureWarning  
 You are setting values through chained assignment. Currently this works in ce  
 A typical example is when you are setting values in a column of a DataFrame,

```
df["col"][row_indexer] = value
```

Use `df.loc[row\_indexer, "col"] = values` instead, to perform the assignment

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_13572\2129877948.py:1: SettingWithCo  
 A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/s>

```
df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan
```

CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax
0	1.0	20-25	4	Ibis	veg	1300.0
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0
2	3.0	25-30	6	RedFox	Veg	1322.0
3	4.0	20-25	-1	LemonTree	Veg	1234.0
4	5.0	35+	3	Ibis	Vegetarian	989.0
5	6.0	35+	3	Ibys	Non-Veg	1909.0
6	7.0	35+	4	RedFox	Vegetarian	1000.0
7	8.0	20-25	7	LemonTree	Veg	2999.0
8	9.0	25-30	2	Ibis	Non-Veg	3456.0
9	10.0	30-35	5	RedFox	non-Veg	NaN

```
df.Age_Group.unique()
```

```
array(['20-25', '30-35', '25-30', '35+'], dtype=object)
```

```
df.Hotel.unique()
```

```
array(['Ibis', 'LemonTree', 'RedFox', 'Ibys'], dtype=object)
```

```
df.Hotel.replace(['Ibys'],'Ibis')  
df.FoodPreference.unique
```

```
<bound method Series.unique of 0      Veg  
1    Non-Veg  
2      Veg  
3      Veg  
4      Veg  
5    Non-Veg  
6      Veg  
7      Veg  
8    Non-Veg  
9    Non-Veg  
Name: FoodPreference, dtype: object>
```

```
df.FoodPreference.replace(['Vegetarian','veg'],'Veg',inplace=True)  
df.FoodPreference.replace(['non-Veg'],'Non-Veg',inplace=True)
```

```
df.EstimatedSalary.fillna(round(df.EstimatedSalary.mean()))  
df.NoOfPax.fillna(round(df.NoOfPax.median()),inplace=True)  
df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()))  
df.Bill.fillna(round(df.Bill.mean()))  
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax
0	1.0	20-25	4	Ibis	Veg	1300.0	
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0	
2	3.0	25-30	6	RedFox	Veg	1322.0	
3	4.0	20-25	-1	LemonTree	Veg	1234.0	
4	5.0	35+	3	Ibis	Veg	989.0	
5	6.0	35+	3	Ibis	Non-Veg	1909.0	
6	7.0	35+	4	RedFox	Veg	1000.0	
7	8.0	20-25	7	LemonTree	Veg	2999.0	
8	9.0	25-30	2	Ibis	Non-Veg	3456.0	
9	10.0	30-35	5	RedFox	Non-Veg	1801.0	

Start coding or generate with AI.





```
# Rannesh Khumar B R
# 240701422
# 14.8.25
# Outliers
```

```
import numpy as np
array=np.random.randint(1,100,16)
def outDetection(array):
    sorted(array)
```

```
def outDetection(array):
    sorted(array)
    Q1,Q3=np.percentile(array,[25,75])
    IQR=Q3-Q1
    lr=Q1-(1.5*IQR)
    ur=Q3+(1.5*IQR)
    return lr,ur
lr,ur=outDetection(array)
```

```
lr,ur
```

```
(np.float64(-15.875), np.float64(121.125))
```

```
import seaborn as sns
import matplotlib.pyplot as plt
sns.distplot(array)
plt.show()
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_6860\579271414.py:3: UserWarning:
```

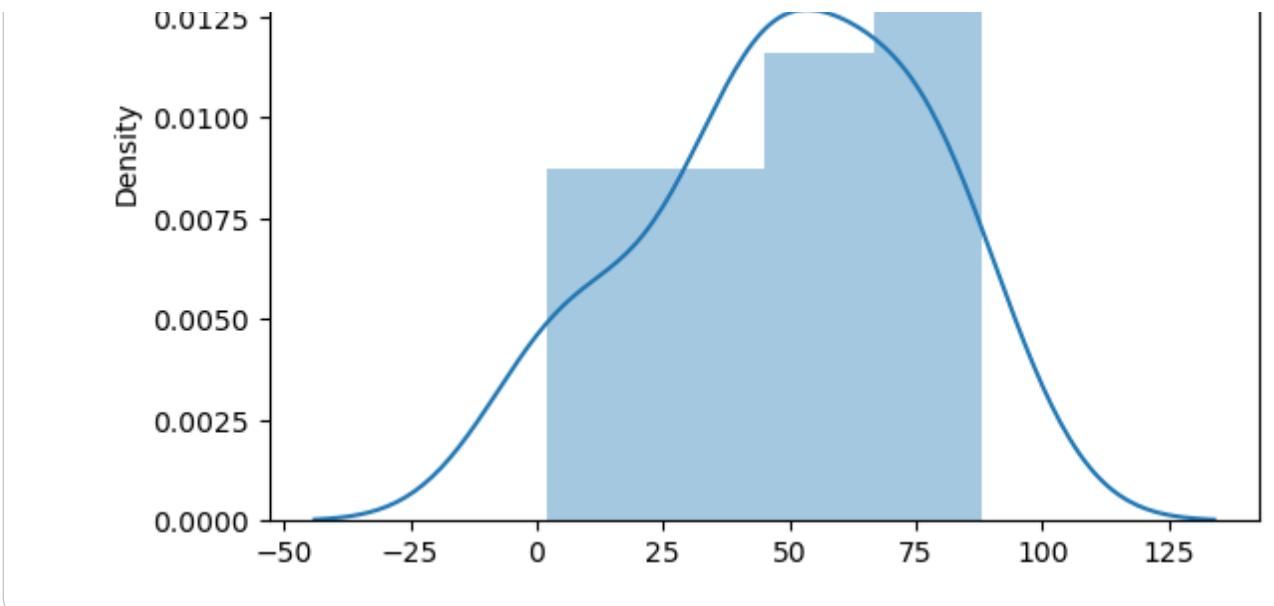
```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

```
sns.distplot(array)
```





Start coding or generate with AI.

Start coding or generate with AI.



```
# Rannesh Khumar B R  
# 240701422  
# 21.8.25  
# Feature scaling
```

```
import numpy as np  
import pandas as pd  
df=pd.read_csv('pre_process_datasample(1).csv')
```

```
df
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes
5	France	35.0	58000.0	Yes
6	Spain	NaN	52000.0	No
7	France	48.0	79000.0	Yes
8	Germany	50.0	83000.0	No
9	France	37.0	67000.0	Yes

```
df.Country.fillna(df.Country.mode()[0])  
features=df.iloc[:, :-1].values
```

```
label=df.iloc[:, -1].values
```

```
from sklearn.impute import SimpleImputer  
age=SimpleImputer(strategy="mean",missing_values=np.nan)  
Salary=SimpleImputer(strategy="mean",missing_values=np.nan)
```

```
age.fit(features[:, [1]])
```

```
▼ SimpleImputer i ?  
SimpleImputer()
```

```
Salary.fit(features[:,[2]])
```

```
  ▾ SimpleImputer ⓘ ⓘ
```

```
SimpleImputer()
```

```
SimpleImputer()
```

```
  ▾ SimpleImputer ⓘ ⓘ
```

```
SimpleImputer()
```

```
features[:,[1]]=age.transform(features[:,[1]])
features[:,[2]]=Salary.transform(features[:,[2]])
features
```

```
array([['France', 44.0, 72000.0],
       ['Spain', 27.0, 48000.0],
       ['Germany', 30.0, 54000.0],
       ['Spain', 38.0, 61000.0],
       ['Germany', 40.0, 63777.7777777778],
       ['France', 35.0, 58000.0],
       ['Spain', 38.77777777777778, 52000.0],
       ['France', 48.0, 79000.0],
       ['Germany', 50.0, 83000.0],
       ['France', 37.0, 67000.0]], dtype=object)
```

```
from sklearn.preprocessing import OneHotEncoder
oh = OneHotEncoder(sparse_output=False)
Country=oh.fit_transform(features[:,[0]])
Country
```

```
array([[1., 0., 0.],
       [0., 0., 1.],
       [0., 1., 0.],
       [0., 0., 1.],
       [0., 1., 0.],
       [1., 0., 0.],
       [0., 0., 1.],
       [1., 0., 0.],
       [0., 1., 0.],
       [1., 0., 0.]])
```

```
final_set=np.concatenate((Country,features[:,[1,2]]),axis=1)
final_set
```

```
array([[1.0, 0.0, 0.0, 44.0, 72000.0],
       [0.0, 0.0, 1.0, 27.0, 48000.0],
```

```
[..., ..., ..., ..., ...],
[0.0, 1.0, 0.0, 30.0, 54000.0],
[0.0, 0.0, 1.0, 38.0, 61000.0],
[0.0, 1.0, 0.0, 40.0, 63777.77777777778],
[1.0, 0.0, 0.0, 35.0, 58000.0],
[0.0, 0.0, 1.0, 38.77777777777778, 52000.0],
[1.0, 0.0, 0.0, 48.0, 79000.0],
[0.0, 1.0, 0.0, 50.0, 83000.0],
[1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)
```

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
sc.fit(final_set)
feat_standard_scaler=sc.transform(final_set)
feat_standard_scaler
```

```
array([[ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
       7.58874362e-01,  7.49473254e-01],
      [-8.16496581e-01, -6.54653671e-01,  1.52752523e+00,
       -1.71150388e+00, -1.43817841e+00],
      [-8.16496581e-01,  1.52752523e+00, -6.54653671e-01,
       -1.27555478e+00, -8.91265492e-01],
      [-8.16496581e-01, -6.54653671e-01,  1.52752523e+00,
       -1.13023841e-01, -2.53200424e-01],
      [-8.16496581e-01,  1.52752523e+00, -6.54653671e-01,
       1.77608893e-01,  6.63219199e-16],
      [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
       -5.48972942e-01, -5.26656882e-01],
      [-8.16496581e-01, -6.54653671e-01,  1.52752523e+00,
       0.00000000e+00, -1.07356980e+00],
      [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
       1.34013983e+00,  1.38753832e+00],
      [-8.16496581e-01,  1.52752523e+00, -6.54653671e-01,
       1.63077256e+00,  1.75214693e+00],
      [ 1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
       -2.58340208e-01,  2.93712492e-01]])
```

```
from sklearn.preprocessing import MinMaxScaler
mms=MinMaxScaler(feature_range=(0,1))
mms.fit(final_set)
feat_minmax_scaler=mms.transform(final_set)
feat_minmax_scaler
```

```
array([[1.          , 0.          , 0.          , 0.73913043, 0.68571429],
      [0.          , 0.          , 1.          , 0.          , 0.          ],
      [0.          , 1.          , 0.          , 0.13043478, 0.17142857],
      [0.          , 0.          , 1.          , 0.47826087, 0.37142857],
      [0.          , 1.          , 0.          , 0.56521739, 0.45079365],
      [1.          , 0.          , 0.          , 0.34782609, 0.28571429],
      [0.          , 0.          , 1.          , 0.51207729, 0.11428571],
      [1.          , 0.          , 0.          , 0.91304348, 0.88571429],
      [0.          , 1.          , 0.          , 1.          , 1.          ],
      [1.          , 0.          , 0.          , 0.43478261, 0.54285714]])
```

Start coding or generate with AI.



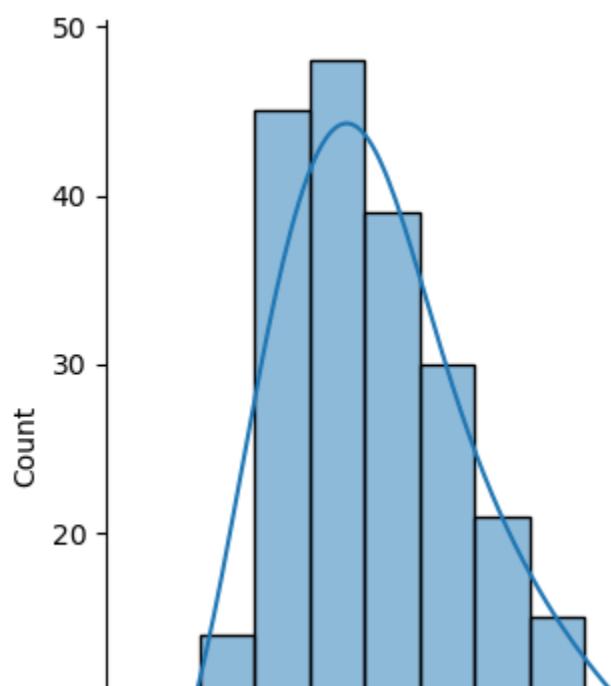
```
# Rannesh Khumar B R  
# 240701422  
# 28.8.25  
# EDA
```

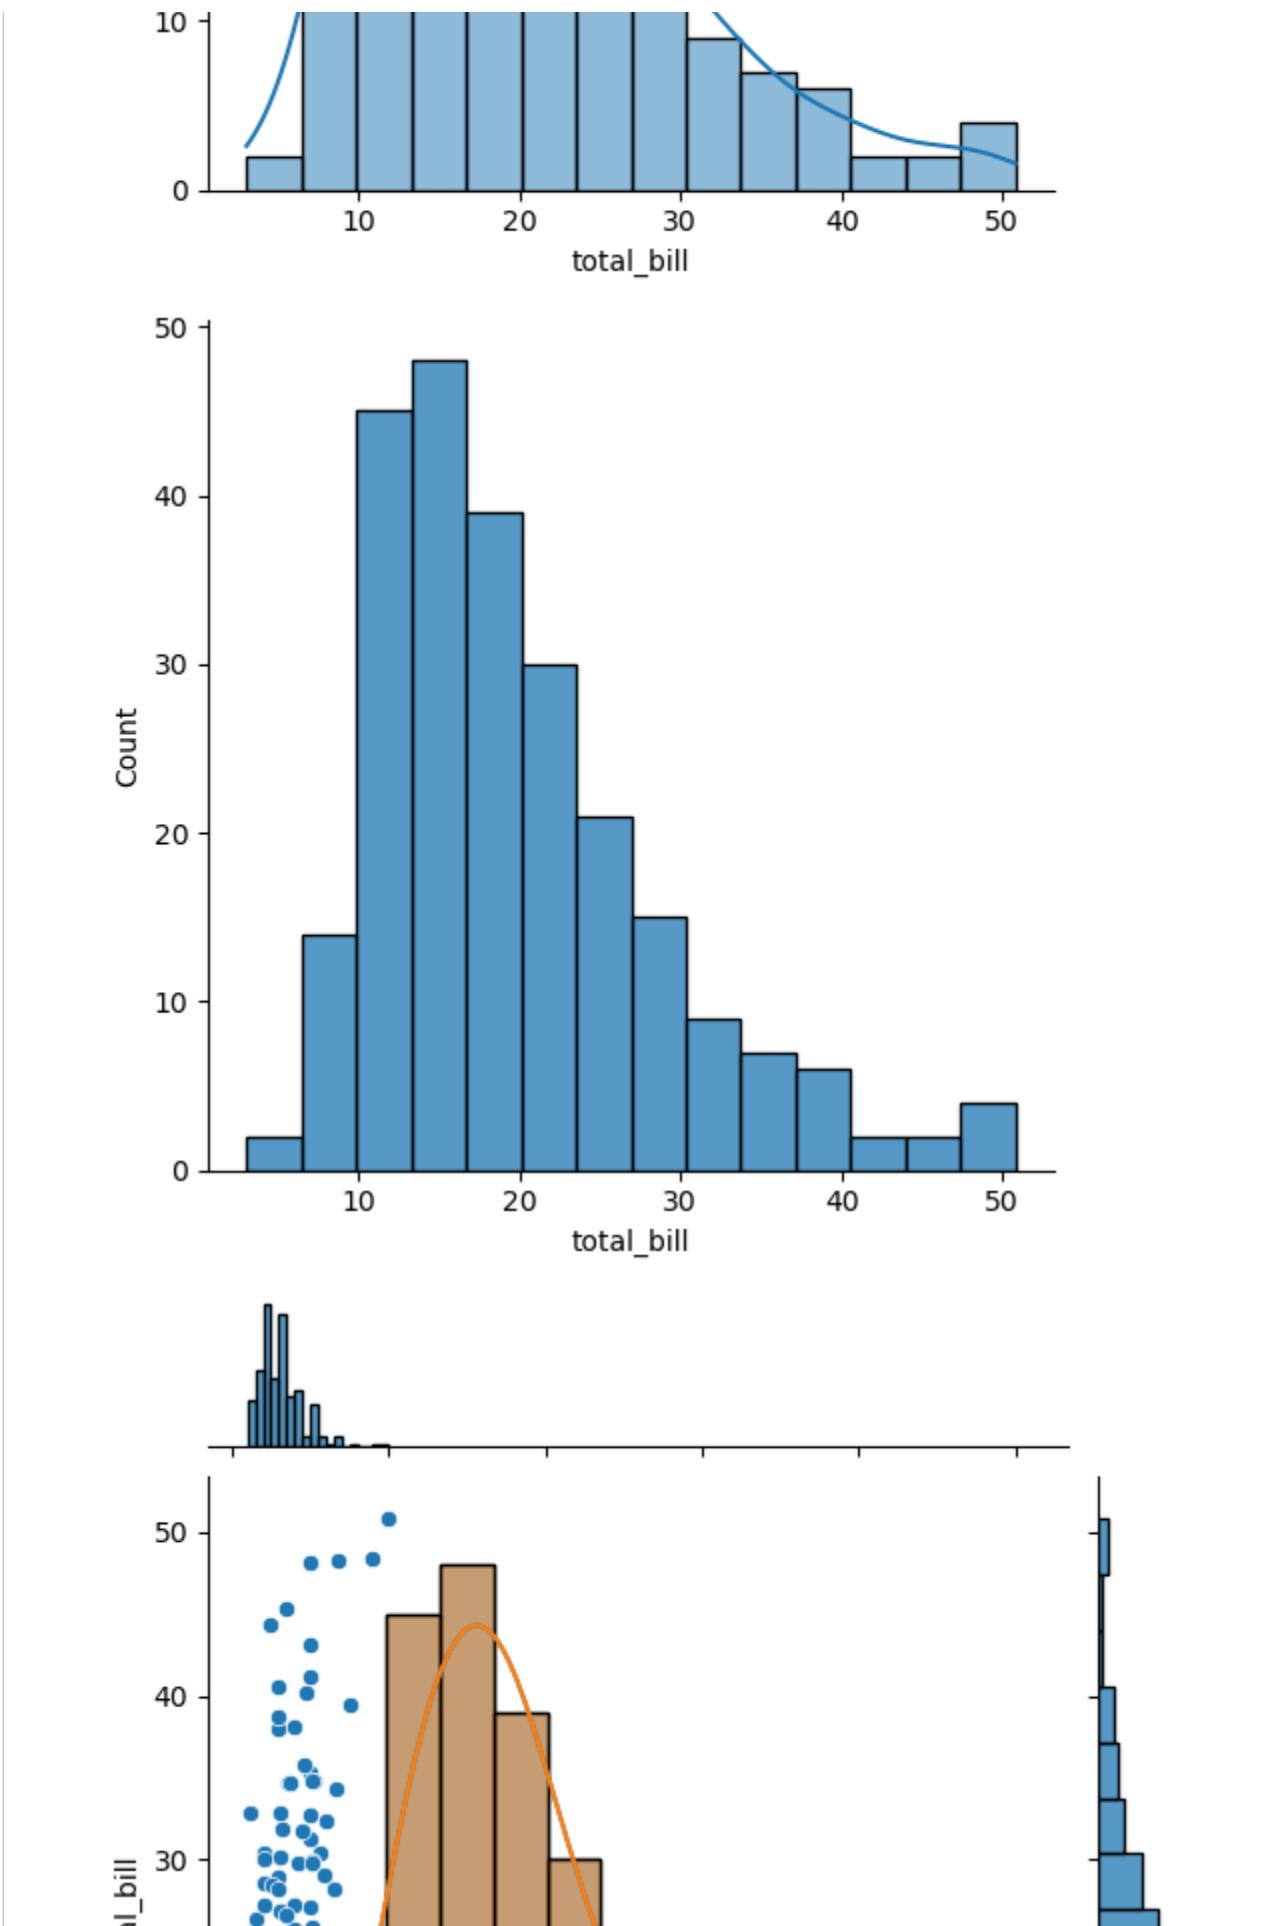
```
import seaborn as sns  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
tips=sns.load_dataset('tips')
```

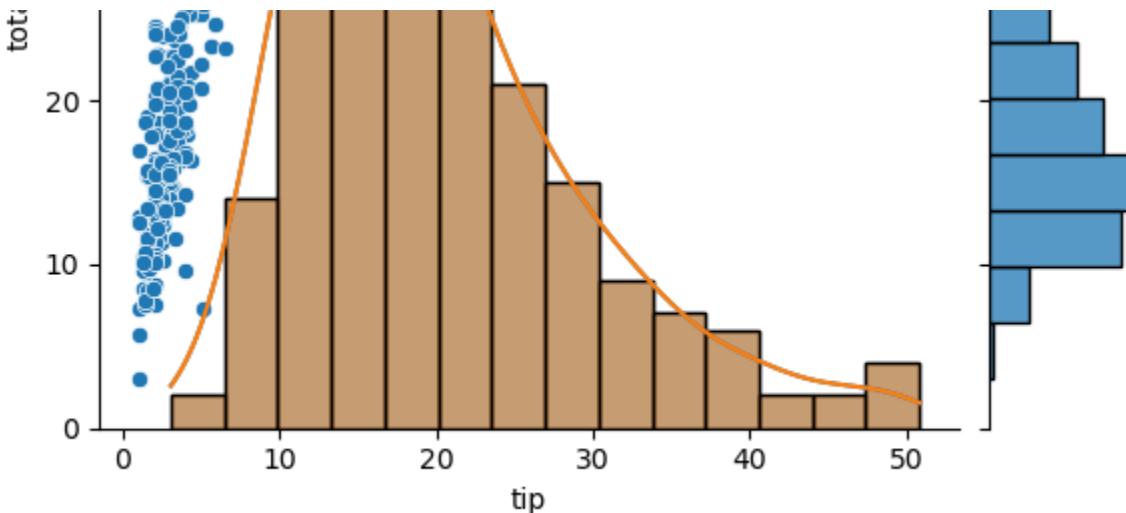
```
tips.head()
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

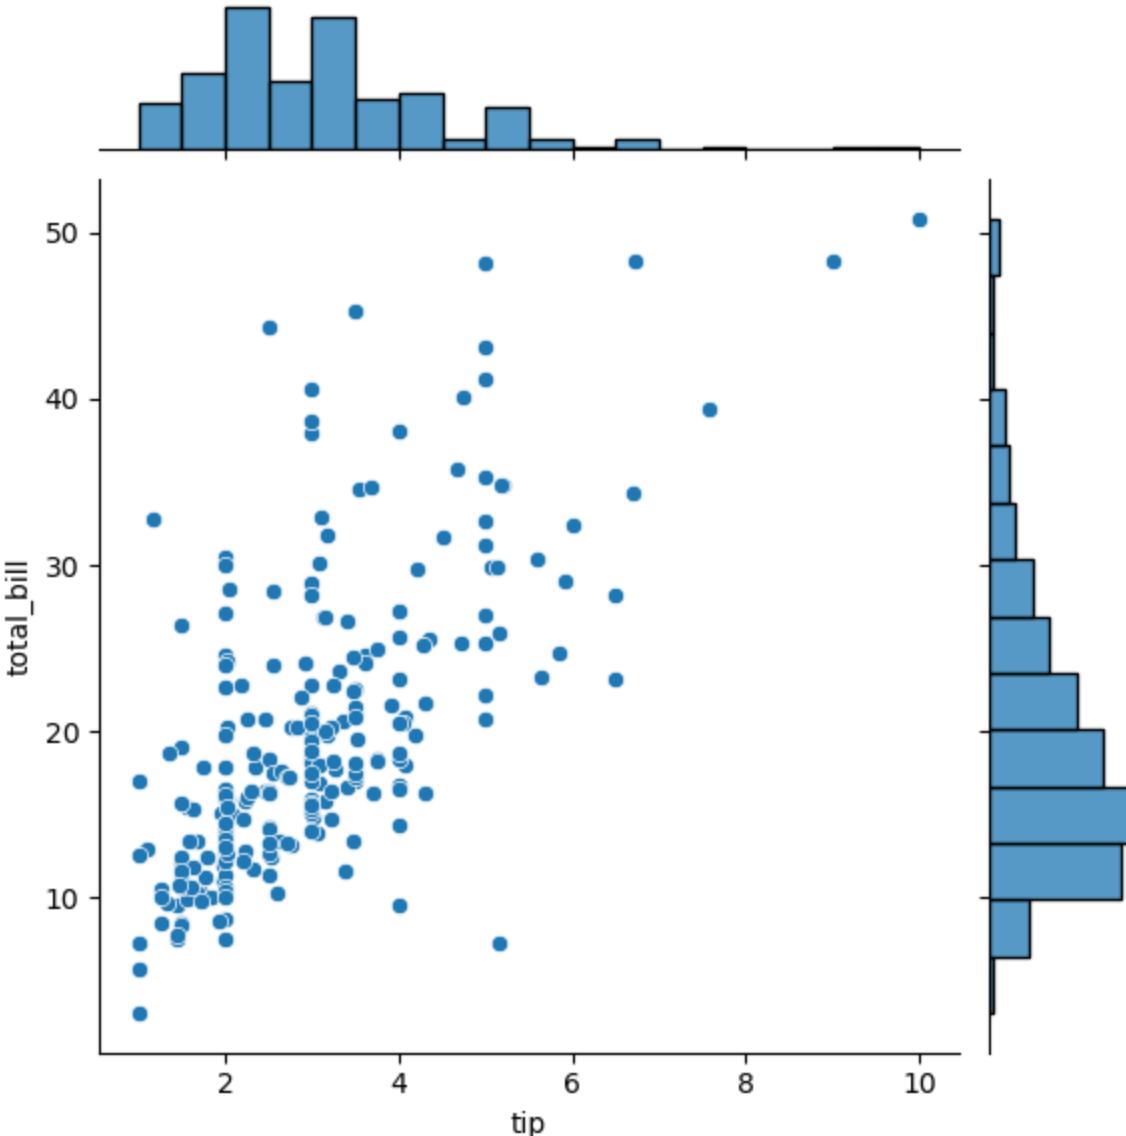
```
sns.histplot(tips.total_bill,kde=True)  
plt.show()
```



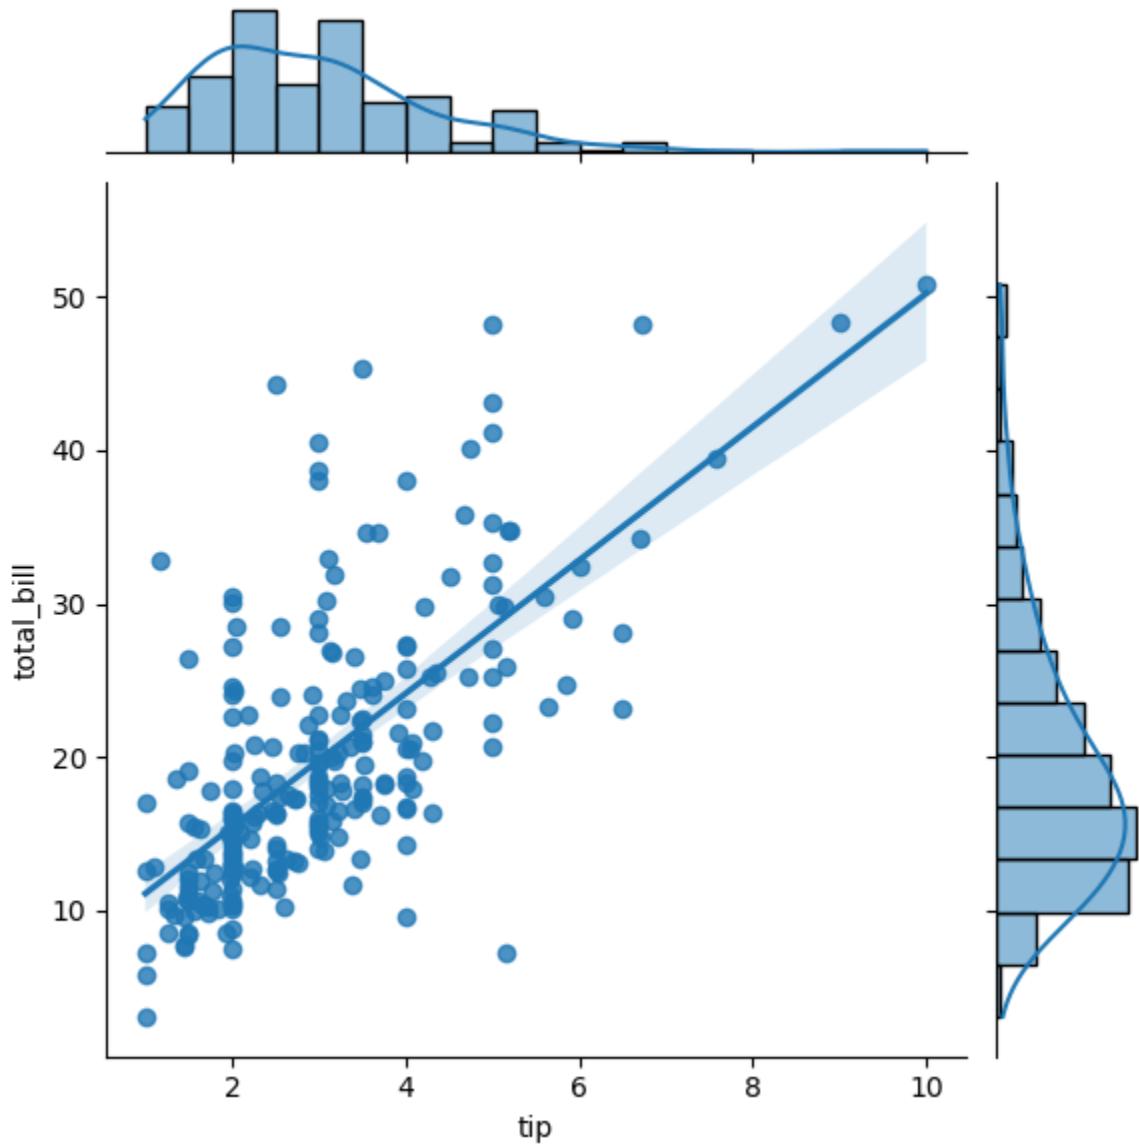




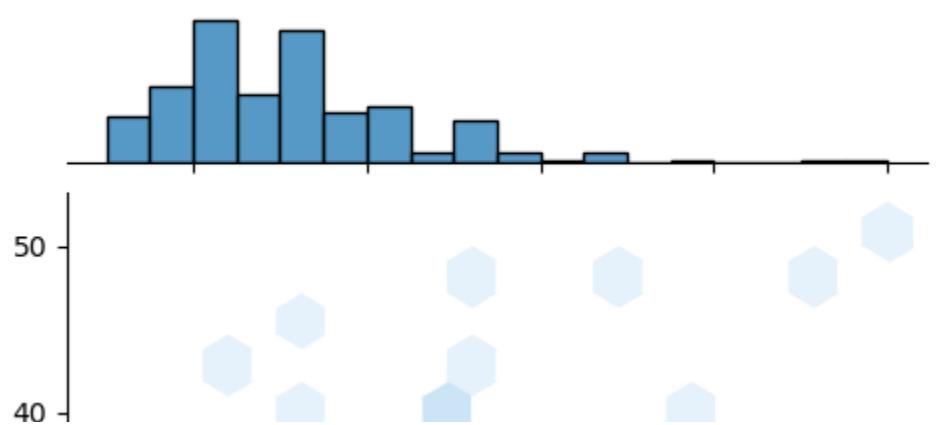
```
sns.jointplot(x=tips.tip,y=tips.total_bill)  
plt.show()
```



```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="reg")
plt.show()
```



```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="hex")
plt.show()
```



```
sns.pairplot(tips)  
plt.show()
```



```
tips.time.value_counts()
```

```
time
Dinner    176
Lunch     68
Name: count, dtype: int64
```

```
sns.pairplot(tips,hue='time')
plt.show()
```

```
sns.pairplot(tips,hue='day')
plt.show()
```

```
sns.heatmap(tips.corr(numeric_only=True), annot=True)  
plt.show()
```

```
sns.boxplot(tips.total_bill)  
plt.show()
```

```
sns.boxplot(tips.tip)  
plt.show()
```

```
sns.countplot(tips.day)
plt.show()
```

```
sns.countplot(tips.sex)
plt.show()
```

```
tips.sex.value_counts().plot(kind='pie')
plt.show()
```

```
tips.sex.value_counts().plot(kind='bar')
plt.show()
```

```
sns.countplot(tips[tips.time=='Dinner']['day'])  
plt.show()
```

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.



```
# Rannesh Khumar B R
# 240701422
# 17.9.25
# Linear Regression
```

```
import numpy as np
import pandas as pd
df=pd.read_csv('Salary_data.csv')
df
```

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891
5	2.9	56642
6	3.0	60150
7	3.2	54445
8	3.2	64445
9	3.7	57189
10	3.9	63218
11	4.0	55794
12	4.0	56957
13	4.1	57081
14	4.5	61111
15	4.9	67938
16	5.1	66029
17	5.3	83088
18	5.9	81363
19	6.0	93940
20	6.8	91738
21	7.1	98273

```
22          7.9  101302
23          8.2  113812
24          8.7  109431
25          9.0  105582
26          9.5  116969
27          9.6  112635
28         10.3  122391
29         10.5  121872
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   YearsExperience  30 non-null      float64 
 1   Salary            30 non-null      int64   
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes
```

```
df.dropna(inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   YearsExperience  30 non-null      float64 
 1   Salary            30 non-null      int64   
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes
```

```
df.describe()
```

	YearsExperience	Salary
<b>count</b>	30.000000	30.000000
<b>mean</b>	5.313333	76003.000000
<b>std</b>	2.837888	27414.429785

	min	1.100000	37731.000000
<b>25%</b>	3.200000	56720.750000	
<b>50%</b>	4.700000	65237.000000	
<b>75%</b>	7.700000	100544.750000	
<b>max</b>	10.500000	122391.000000	

```
features=df.iloc[:,[0]].values  
label=df.iloc[:,[1]].values
```

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_state=42)
```

```
from sklearn.linear_model import LinearRegression  
model=LinearRegression()  
model.fit(x_train,y_train)
```

```
▼ LinearRegression ⓘ ⓘ  
LinearRegression()
```

```
model.score(x_train,y_train)
```

```
0.9577907749872991
```

```
model.score(x_test,y_test)
```

```
0.9531732818427658
```

```
model.coef_
```

```
array([[9339.90339715]])
```

```
model.intercept_
```

```
array([26561.50676243])
```

```
import pickle  
pickle.dump(model,open('SalaryPred.model','wb'))
```

```
model=pickle.load(open('SalaryPred.model','rb'))
```

```
yr_of_exp=float(input("Enter Years of Experience: "))
```

```
y=yr_of_exp  
yr_of_exp_NP=np.array([[yr_of_exp]])  
Salary=model.predict(yr_of_exp_NP)  
  
print("Estimated salary for {} year of exp is{}".format(yr_of_exp,Salary))
```

```
Enter Years of Experience: 33  
estimated salary for 33.0 year of exp is[[334778.31886853]]
```

```
array([[ 39343],  
       [ 46205],  
       [ 37731],  
       [ 43525],  
       [ 39891],  
       [ 56642],  
       [ 60150],  
       [ 54445],  
       [ 64445],  
       [ 57189],  
       [ 63218],  
       [ 55794],  
       [ 56957],  
       [ 57081],  
       [ 61111],  
       [ 67938],  
       [ 66029],  
       [ 83088],  
       [ 81363],  
       [ 93940],  
       [ 91738],  
       [ 98273],  
       [101302],  
       [113812],  
       [109431],  
       [105582],  
       [116969],  
       [112635],  
       [122391],  
       [121872]])
```

Start coding or generate with AI.





```
# Rannesh Khumar B R  
# 240701422  
# 17.9.25  
# Logistic Regression
```

```
import numpy as np  
import pandas as pd  
df=pd.read_csv('Social_Network_Ads.csv')  
df
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

```
features=df.iloc[:,[2,3]].values  
label=df.iloc[:,4].values  
features
```

```
array([[ 19, 19000],  
       [ 35, 20000],  
       [ 26, 43000],  
       [ 27, 57000],  
       [ 19, 76000],  
       [ 27, 58000],  
       [ 27, 84000],  
       [ 32, 150000],  
       [ 25, 33000],  
       [ 35, 65000],  
       [ 26, 80000],  
       [ 26, 52000]])
```

```
L      20,  22000],
[ 20,  86000],
[ 32,  18000],
[ 18,  82000],
[ 29,  80000],
[ 47,  25000],
[ 45,  26000],
[ 46,  28000],
[ 48,  29000],
[ 45,  22000],
[ 47,  49000],
[ 48,  41000],
[ 45,  22000],
[ 46,  23000],
[ 47,  20000],
[ 49,  28000],
[ 47,  30000],
[ 29,  43000],
[ 31,  18000],
[ 31,  74000],
[ 27,  137000],
[ 21,  16000],
[ 28,  44000],
[ 27,  90000],
[ 35,  27000],
[ 33,  28000],
[ 30,  49000],
[ 26,  72000],
[ 27,  31000],
[ 27,  17000],
[ 33,  51000],
[ 35,  108000],
[ 30,  15000],
[ 28,  84000],
[ 23,  20000],
[ 25,  79000],
[ 27,  54000],
[ 30,  135000],
[ 31,  89000],
[ 24,  32000],
[ 18,  44000],
[ 29,  83000],
[ 35,  23000],
[ 27,  58000],
[ 24,  55000],
[ 23,  48000],
[ 28,  79000],
```

### label

```
array([0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
```

```
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression
```

```
for i in range(1,401):
    x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2)
    model=LogisticRegression()
    model.fit(x_train,y_train)
    train_score=model.score(x_train,y_train)
    test_score=model.score(x_test,y_test)
```

```
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,  
finalModel=LogisticRegression()  
finalModel.fit(x_train,y_train)
```

▼ LogisticRegression [i](#) [?](#)  
LogisticRegression()

```
print(finalModel.score(x_train,y_train))  
print(finalModel.score(x_test,y_test))
```

0.846875  
0.8

```
from sklearn.metrics import classification_report  
print(classification_report(label,finalModel.predict(features)))
```

	precision	recall	f1-score	support
0	0.84	0.92	0.88	257
1	0.82	0.69	0.75	143
accuracy			0.84	400
macro avg	0.83	0.81	0.82	400

weighted avg	0.84	0.84	0.83	400
--------------	------	------	------	-----

Start coding or generate with AI.



```
# Rannesh Khumar B R  
# 240701422  
# 24.10.25  
# K means Clustering
```

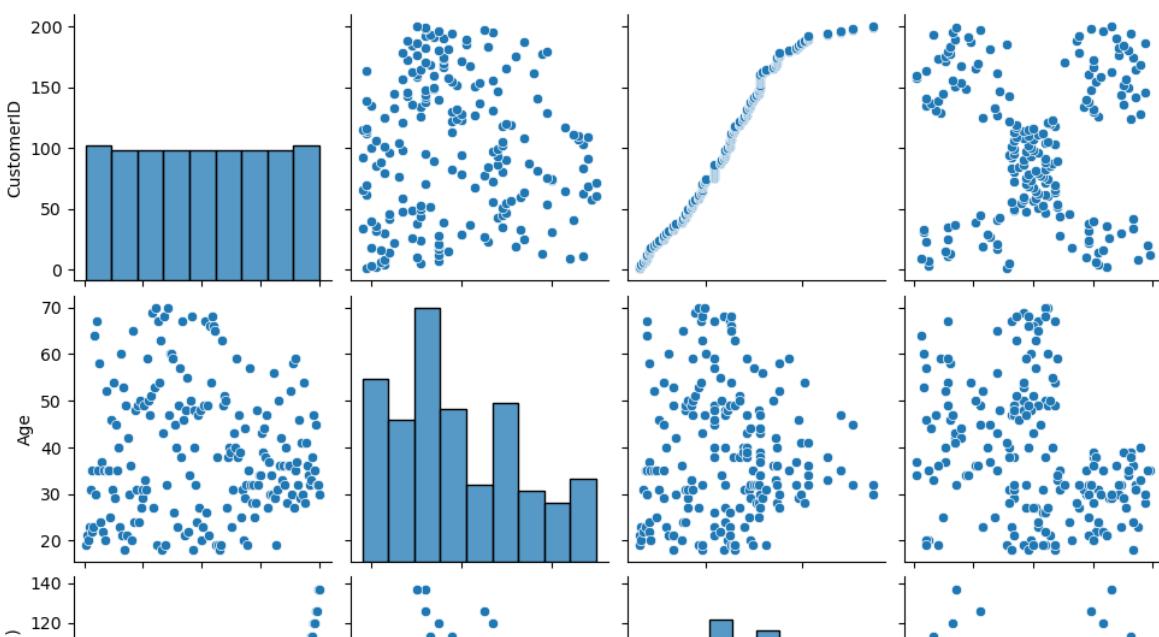
```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

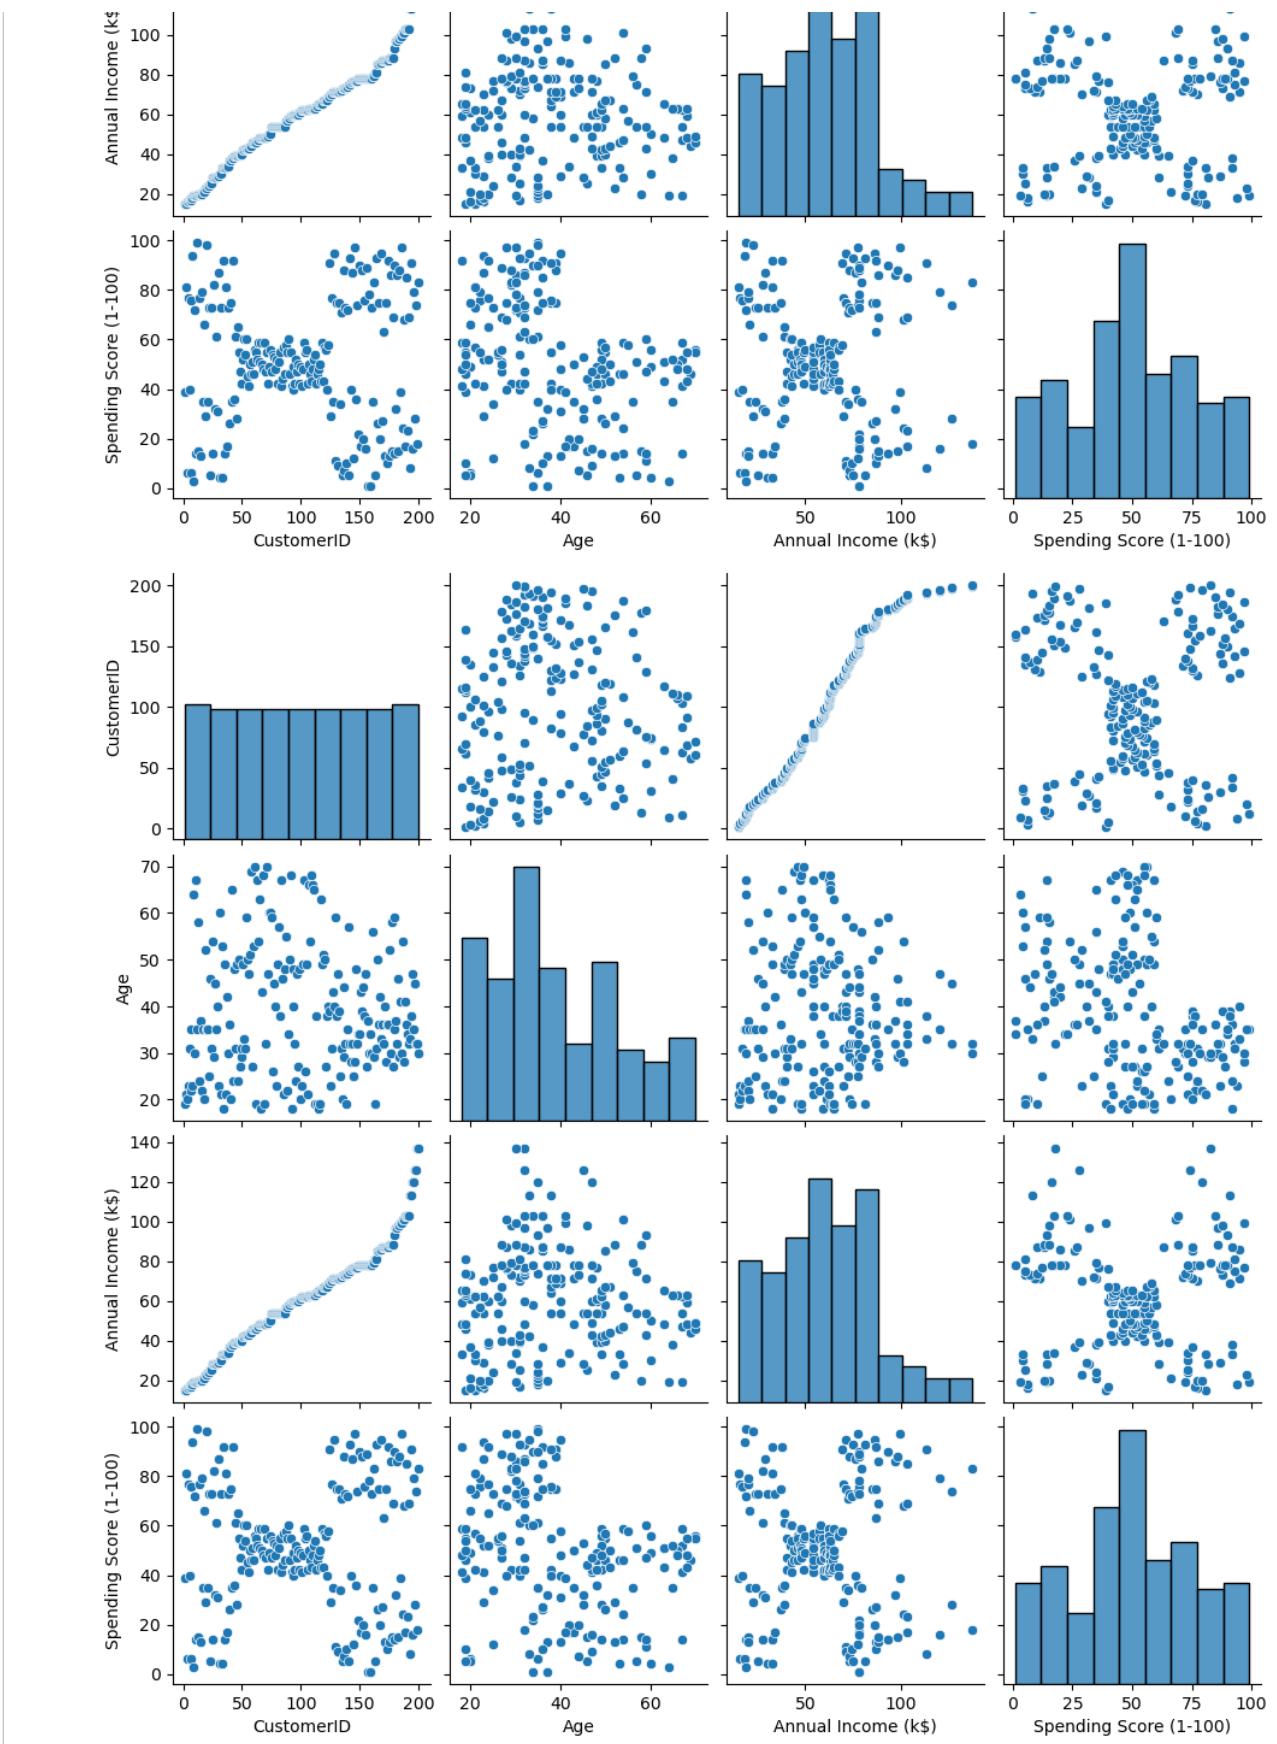
```
df=pd.read_csv('Mall_Customers.csv')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 5 columns):  
 #   Column           Non-Null Count  Dtype    
 ---  --    
 0   CustomerID      200 non-null    int64  
 1   Gender          200 non-null    object  
 2   Age             200 non-null    int64  
 3   Annual Income (k$) 200 non-null    int64  
 4   Spending Score (1-100) 200 non-null    int64  
dtypes: int64(4), object(1)  
memory usage: 7.9+ KB
```

```
sns.pairplot(df)  
plt.show()
```





```
features=df.iloc[:,[3,4]].values
```

```
from sklearn.cluster import KMeans  
model=KMeans(n_clusters=5)  
model.fit(features)  
KMeans(n_clusters=5)
```

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: U  
warnings.warn(  
    ▾ KMeans ⓘ ⓘ  
    KMeans(n_clusters=5)
```

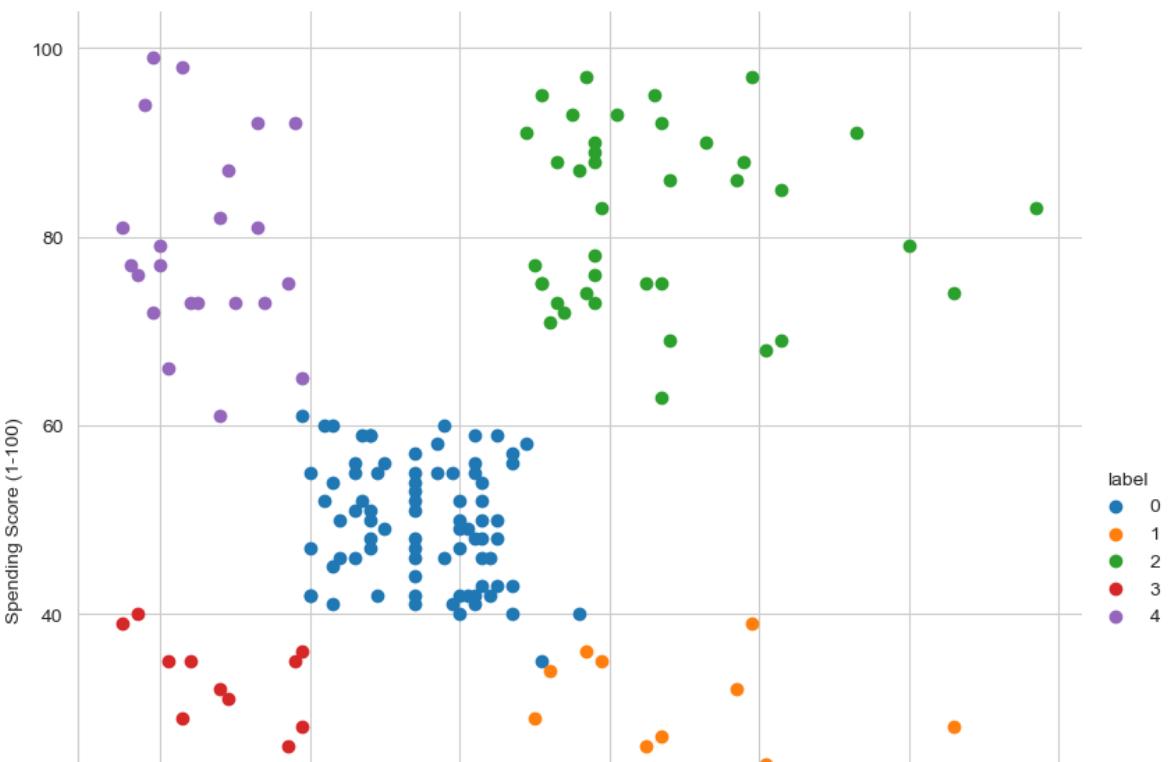
```
Final=df.iloc[:,[3,4]]  
Final['label']=model.predict(features)  
Final.head()
```

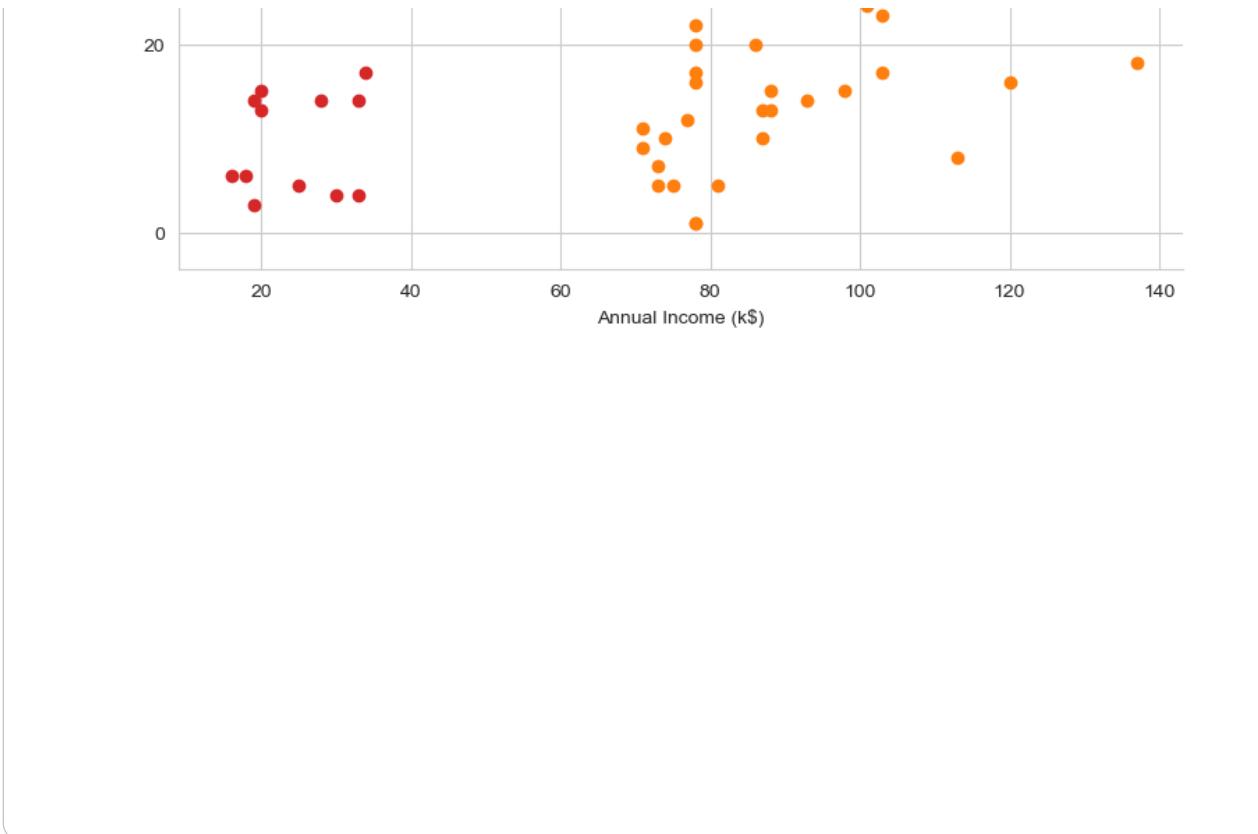
C:\Users\HP\AppData\Local\Temp\ipykernel\_15280\470183701.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.loc.html>

	Annual Income (k\$)	Spending Score (1-100)	label
0	15	39	3
1	15	81	4
2	16	6	3
3	16	77	4
4	17	40	3

```
sns.set_style("whitegrid")  
sns.FacetGrid(Final,hue="label",height=8) \  
.map(plt.scatter,"Annual Income (k$)", "Spending Score (1-100)") \  
.add_legend();  
plt.show()
```

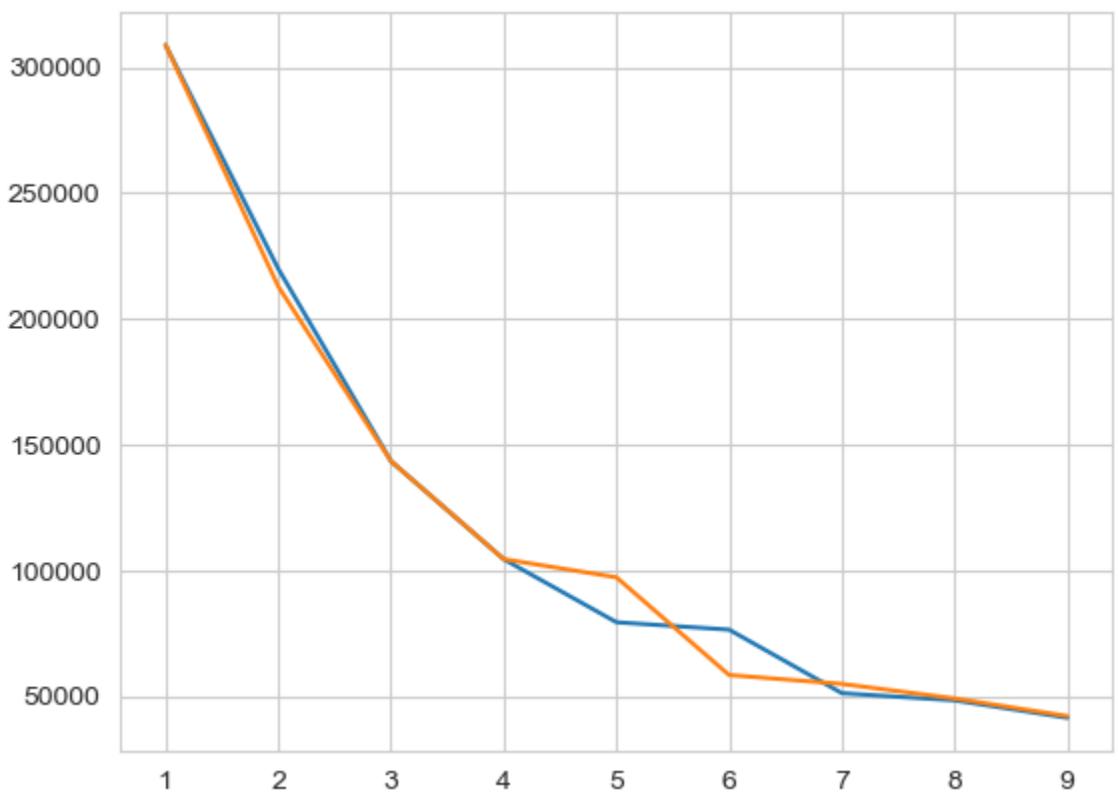




```
features_el=df.iloc[:,[2,3,4]].values
from sklearn.cluster import KMeans
wcss=[]
for i in range(1,10):
    model=KMeans(n_clusters=i)
    model.fit(features_el)
    wcss.append(model.inertia_)
plt.plot(range(1,10),wcss)
```

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: U
  warnings.warn(
[<matplotlib.lines.Line2D at 0x246b722c410>]
```

```
plt.show()
```



Start coding or generate with AI.



```
# Rannesh Khumar B R  
# 240701422  
# 1.10.25  
# KNN
```

```
import numpy as np  
import pandas as pd
```

```
df=pd.read_csv('Iris.csv')  
df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 150 entries, 0 to 149  
Data columns (total 6 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --     
 0   Id               150 non-null    int64    
 1   SepalLengthCm   150 non-null    float64  
 2   SepalWidthCm    150 non-null    float64  
 3   PetalLengthCm   150 non-null    float64  
 4   PetalWidthCm    150 non-null    float64  
 5   Species          150 non-null    object    
 dtypes: float64(4), int64(1), object(1)  
memory usage: 7.2+ KB
```

```
df.Species.value_counts()  
df
```

	<b>ID</b>	<b>SepalLengthCm</b>	<b>SepalWidthCm</b>	<b>PetalLengthCm</b>	<b>PetalWidthCm</b>	<b>Species</b>
<b>0</b>	1	5.1	3.5	1.4	0.2	Iris-setosa
<b>1</b>	2	4.9	3.0	1.4	0.2	Iris-setosa
<b>2</b>	3	4.7	3.2	1.3	0.2	Iris-setosa
<b>3</b>	4	4.6	3.1	1.5	0.2	Iris-setosa
<b>4</b>	5	5.0	3.6	1.4	0.2	Iris-setosa
...	...	...	...	...	...	...
<b>145</b>	146	6.7	3.0	5.2	2.3	Iris-virginica
<b>146</b>	147	6.3	2.5	5.0	1.9	Iris-virginica

View in Colab

```
features=df.iloc[:, :-1].values  
label=df.iloc[:, 5].values
```

```
from sklearn.model_selection import train_test_split  
from sklearn.neighbors import KNeighborsClassifier
```

```
xtrain,xtest,ytrain,ytest=train_test_split(features,label,test_size=.2,random_state=42)  
model_KNN=KNeighborsClassifier(n_neighbors=5)  
model_KNN.fit(xtrain,ytrain)
```

▼ KNeighborsClassifier ⓘ ⓘ

```
KNeighborsClassifier()
```

```
print(model_KNN.score(xtrain,ytrain))  
print(model_KNN.score(xtest,ytest))
```

```
1.0
```

```
1.0
```

```
from sklearn.metrics import confusion_matrix  
confusion_matrix(label,model_KNN.predict(features))
```

```
array([[50,  0,  0],  
       [ 0, 50,  0],  
       [ 0,  0, 50]])
```

```
from sklearn.metrics import classification_report  
print(classification_report(label,model_KNN.predict(features)))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	50
Iris-versicolor	1.00	1.00	1.00	50
Iris-virginica	1.00	1.00	1.00	50
accuracy			1.00	150
macro avg	1.00	1.00	1.00	150
weighted avg	1.00	1.00	1.00	150





```
#Rannesh Khumar B R  
#240701422  
# 8.10.25  
#T-test
```

```
import numpy as np  
from scipy import stats  
  
marks = np.array([72, 68, 75, 70, 74, 69, 71, 73, 70, 72])  
  
mu_0 = 70  
  
t_stat, p_value = stats.ttest_1samp(marks, mu_0)  
print(f"T-statistic: {t_stat:.3f}")  
print(f'P-value: {p_value:.4f}')  
alpha = 0.05  
if p_value < alpha:  
    print('Reject Null Hypothesis → Mean is significantly different from 70')  
else:  
    print('Fail to Reject Null Hypothesis → No significant difference.')
```

```
T-statistic: 1.993  
P-value: 0.0774  
Fail to Reject Null Hypothesis → No significant difference.
```

Start coding or generate with AI.





```
# Rannesh Khumar B R
# 240701422
# 8.10.25
# Z-test
```

```
import numpy as np
from math import sqrt
from scipy.stats import norm
# Given data
x_bar = 51.2 # sample mean
mu_0 = 50 # population mean
sigma = 3 # population standard deviation
n = 36 # sample size
# Calculate Z-statistic
z_stat = (x_bar - mu_0) / (sigma / sqrt(n))
# Two-tailed p-value
p_value = 2 * (1 - norm.cdf(abs(z_stat)))
print(f"Z-statistic: {z_stat:.3f}")
print(f"P-value: {p_value:.4f}")
alpha = 0.05
if p_value < alpha:
    print("Reject Null Hypothesis → Mean is significantly different from 50 g.")
else:
    print("Fail to Reject Null Hypothesis → No significant difference.")
```

```
Z-statistic: 2.400
P-value: 0.0164
Reject Null Hypothesis → Mean is significantly different from 50 g.
```





```
# Rannesh Khumar B R
# 240701422
# 8.10.25
# Anova test
```

```
import numpy as np
from scipy import stats

# Data
A = [20, 22, 23]
B = [19, 20, 18]
C = [25, 27, 26]

# Perform one-way ANOVA
f_stat, p_value = stats.f_oneway(A, B, C)

print(f"F-statistic: {f_stat:.3f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
    print("Reject Null Hypothesis → Means are significantly different.")
else:
    print("Fail to Reject Null Hypothesis → No significant difference.")
```

```
F-statistic: 25.923
P-value: 0.0011
Reject Null Hypothesis → Means are significantly different.
```

